

# Database Systems I

CMPT 354 Summer 2024

Zhengjie Miao

# Course Setup

Ask professor if he can post the due dates of the assignment

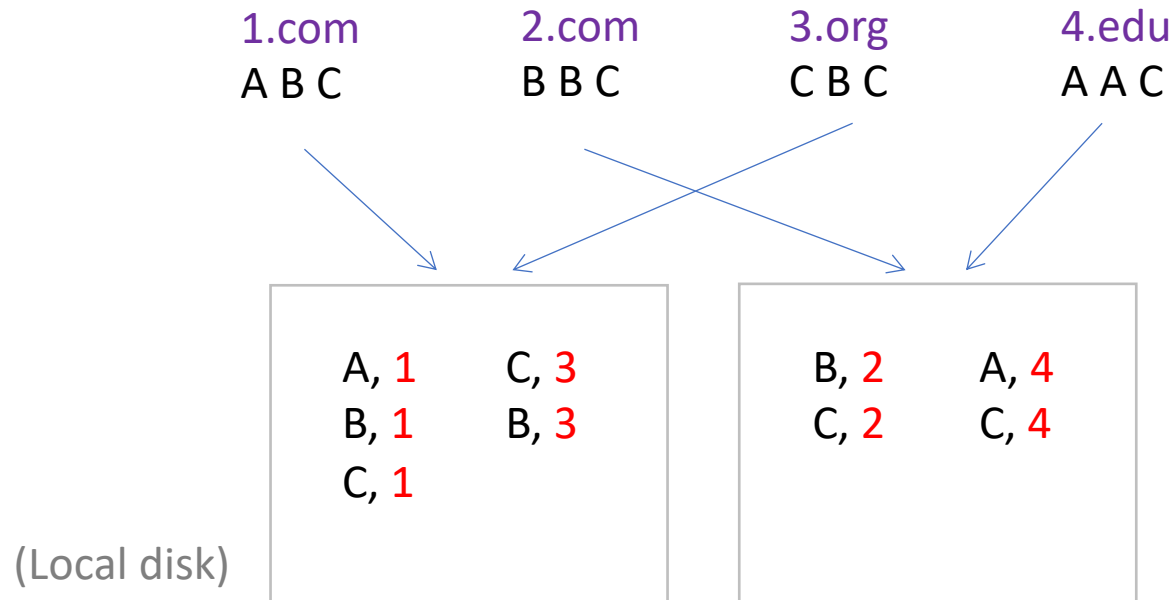
- OH slots:
  - Tue 2pm-3pm, Zhengjie@TASC 1 9407
  - Mon 11am-12pm, Xinyi@ASB9812
  - Wed 1pm-2pm, Obumneme@ASB9812
  - Thu 3pm-4pm, Zahra@ASB9812
- Sign up for Piazza!
  - The place to ask course-related questions
  - <https://piazza.com/sfu.ca/summer2024/cmpt354d100>
  - Access code: qc76kdzsz9l

# Recap: MapReduce

Online analytical processing

- Which design is better for OLAP?
    - MapReduce vs. Relational DBMS?
  - Many problems can be processed in MapReduce pattern:
    - Given a lot of unsorted data Allows you to divide your data among multiple machines
    - **Map**: extract something of interest from each record
    - **Shuffle**: group the intermediate results in some way
    - **Reduce**: further process (e.g., aggregate, summarize, analyze, transform) each group and write final results
- (Customize map and reduce for problem at hand)

# MR Example - Map



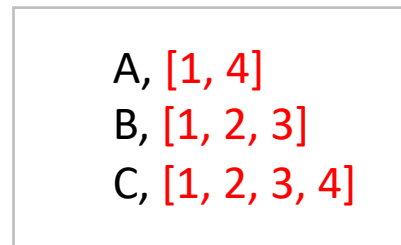
key: document id  
value: document contents

`map(String key, String value):`

for each word `w` in value:

`Emit_Intermediate(w, document id);`

# MR Example - Shuffle & Reduce



Values exchanged by shuffle process

key: a word

values: a list of doc id

Why is MapReduced criticized?

Proved to be inefficient in database systems.

Ask professor why it is inefficient.

The reduce function accepts all pairs for a given word, sorts the corresponding document IDs

MapReduce does not require any schema

# Trends

Typically, you will need to maintain a transactional and analytical system, where you move data between them.

Ask him about why we would move our data frequently?

- 1 Hybrid Transactional and Analytical Processing

Maintain database inside virtual machine

Cloud databases could allow infinite storage

disaggregating the storage and computation layer?

The scale of computation is independent from storage? Ask the professor what he means

- 2. Cloud-Native Databases

- 3. Lakehouse = Data Lake + Data Warehouse

Search up Data Lake and Data Warehouse

Lakehouse wants to result in efficient queries

- 4. Specialized Systems

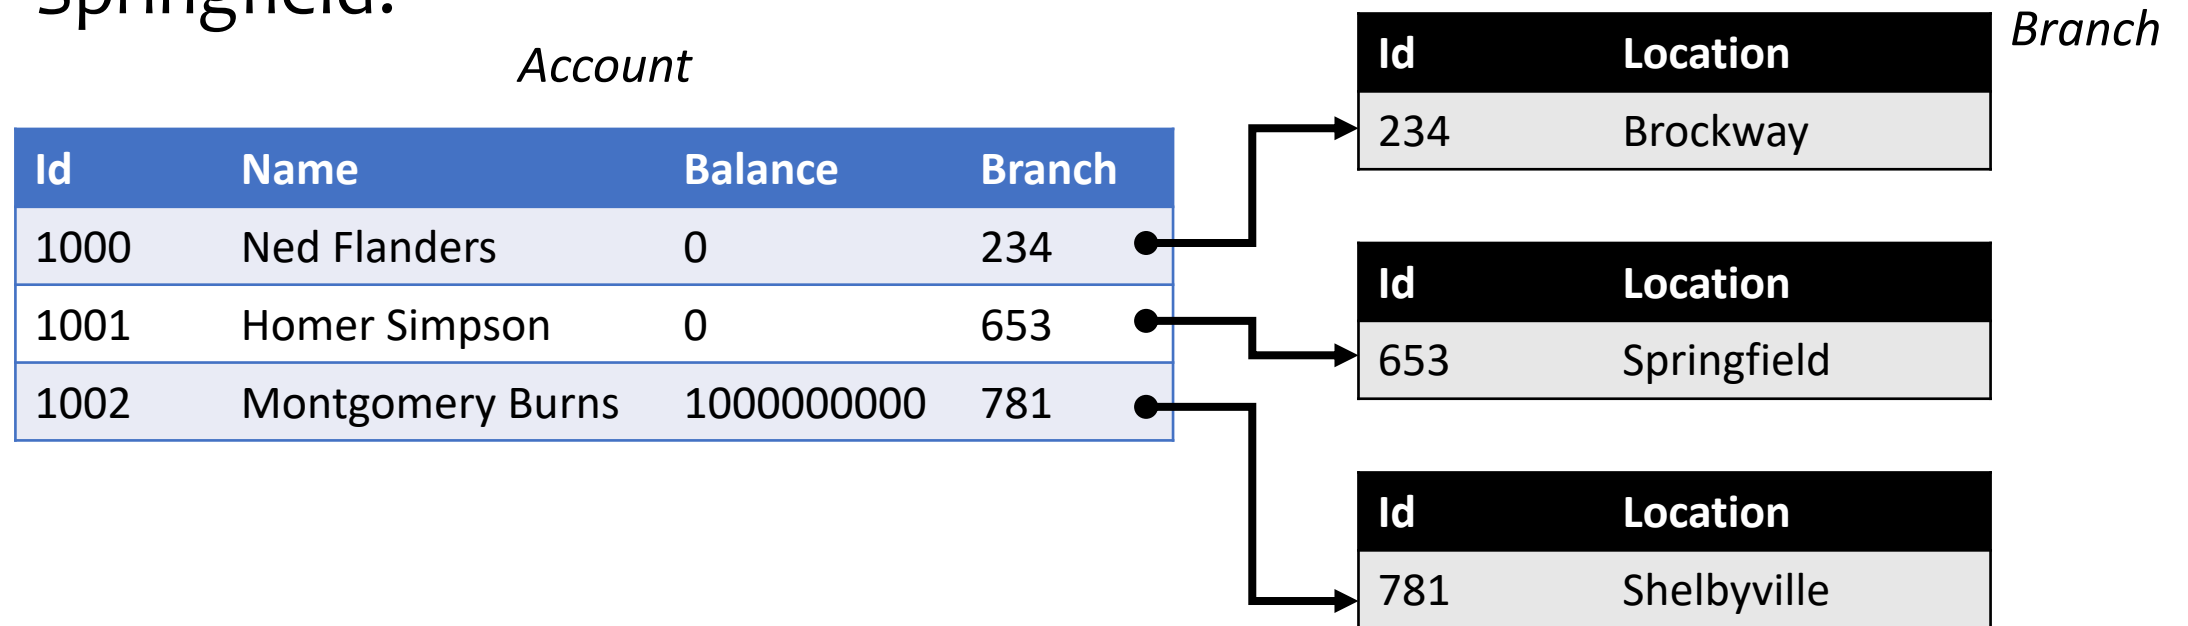
- E.g. Time-series Database, GPU Acceleration, Vector Database

# Outline

- An overview of data models
- Basics of the Relational Model
- Relational Algebra
  - Core operators & Derived operators
  - Extension
  - Relational calculus

# Recap: Data Models

- Query: Who have accounts with 0 balance managed by a branch in Springfield?



- Programmer controls “navigation”: accounts → branches
  - How about branches → accounts? (e.g., find branches managing accounts with 0 balance)



# Data Model

- Data Model
  - mathematical formalism (or conceptual way) for describing the data

If you do not have a data model, what happens?  
Ask professor
- The description generally consists of three parts:
  - Structure of the data
  - Operations on the data
  - Constraints on the data

# Structure of the data

- Schema (e.g., table names, attribute names)
  - Describe the **conceptual** structure of the data
- Different from **data structure** (e.g., list, array)
  - Data structure can be seen as a **physical** data model

# Operations on the data

- Query language (e.g., RA, SQL)
  - Describe what operations that can be performed on data
- Two kinds of operations
  - operations that retrieve information (“query” the data)
  - operations that change the database (“update” the data)
- Different from programming languages (e.g., C, Java)
  - Support a set of limited operations
  - Allow for query optimizations

Your query can be optimized if the system finds another query which is equivalent and more efficient

Programming languages can perform operations which the query language can not

# Constraints on the data.

- Constraints (e.g.,  $\text{balance} \geq 0$ , account id is unique)
  - describe limitations on what the data can be.
- Different kinds of constraints
  - Domain constraints Ex:  $\text{balance} \geq 0$
  - Integrity constraints Ex: the user id is unique
- Why does it matter?
  - Ensure the correctness of data

# Commonly Used Data Models

- Relational Data Model
- Key-Value Data Model
- Semi-structured Data Model (e.g., Json, XML)

# The Relational Model in Brief

Id	Name	Balance	Branch		Id	Location
1000	Ned Flanders	0	234	●	234	Brockway
1001	Homer Simpson	0	653	●	653	Springfield
1002	Montgomery Burns	1000000000	781	●	781	Shelbyville

I thought SQL used RA, but that RA is not the query language

- Structure of the Data
  - Table structure
- Query language
  - RA, RC, SQL
- Constraints on the data
  - E.g., id is unique, balance  $\geq 0$ , name is not NULL

# The Key-Value Model in Brief

Key → Value
1000 → (Ned Flanders, 0, 234)
1001 → (Homer Simpson, 0, 653)
1002 → (Montgomery Burns, 1000000000, 781)

- Structure of the Data
  - (Key, Value) pairs
  - Key is an integer/string, value can be any object
- Query language
  - get(key), put(key, value)
- Constraints on the data
  - E.g., key is unique, value is not NULL

# The Semistructured Model in Brief

You cannot apply the same operation to each account?  
Ask professor to clarify this


- Structure of the Data
  - Tree structure (e.g., XML, json)
- Query language
  - XQuery, MongoDB QL
- Constraints
  - E.g., each <Account> has a <Id> element and a <name> element nested within

Resembles the hierarchical model  
“what goes around comes around”

```
<Accounts>
  <Account>
    <Id>1000</Id>
    <Name>Ned Flanders</Name>
    <Balance>0</Balance>
    <Branch>234</Branch>
  </Account>
  <Account>
    <Id>1001</Id>
    <Name>Homer Simpson</Name>
    <Balance>0</Balance>
    <Branch>653</Branch>
  </Account>
  <Account>
    <Id>1002</Id>
    <Name>Montgomery Burns</Name>
    <Balance>1000000000</Balance>
    <Branch>781</Branch>
  </Account>
</Accounts>
```



# Edgar F. Codd (1923-2003)



- Pilot in the Royal Air Force in WW2
- Inventor of the relational model and algebra while at IBM
- Turing Award, 1981

# Relational data model

- A database is a collection of **relations** (or **tables**)
- Each relation has a set of **attributes** (or **columns**)
- Each attribute has a name and a **domain** (or **type**)
  - Set-valued attributes are not allowed **Attributes must be atomic**
- Each relation contains a set of **tuples** (or **rows**)
  - Each tuple has a value for each attribute of the relation
  - Duplicate tuples are not allowed
    - Two tuples are duplicates if they agree on all attributes **Set semantics**
- Simplicity is a virtue!

# Example

*User*

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...	...	...	...

Ordering of rows doesn't matter  
(even though output is  
always in some order)

*Group*

<i>gid</i>	<i>name</i>
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

*Member*

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

# Schema vs. instance

- Schema (metadata)

- Specifies the logical structure of data
- Is defined at setup time
- Rarely changes

The schema rarely changes since we must update the change to the schema to all instances

- Instance

- Represents the data content
- Changes rapidly, but always conforms to the schema

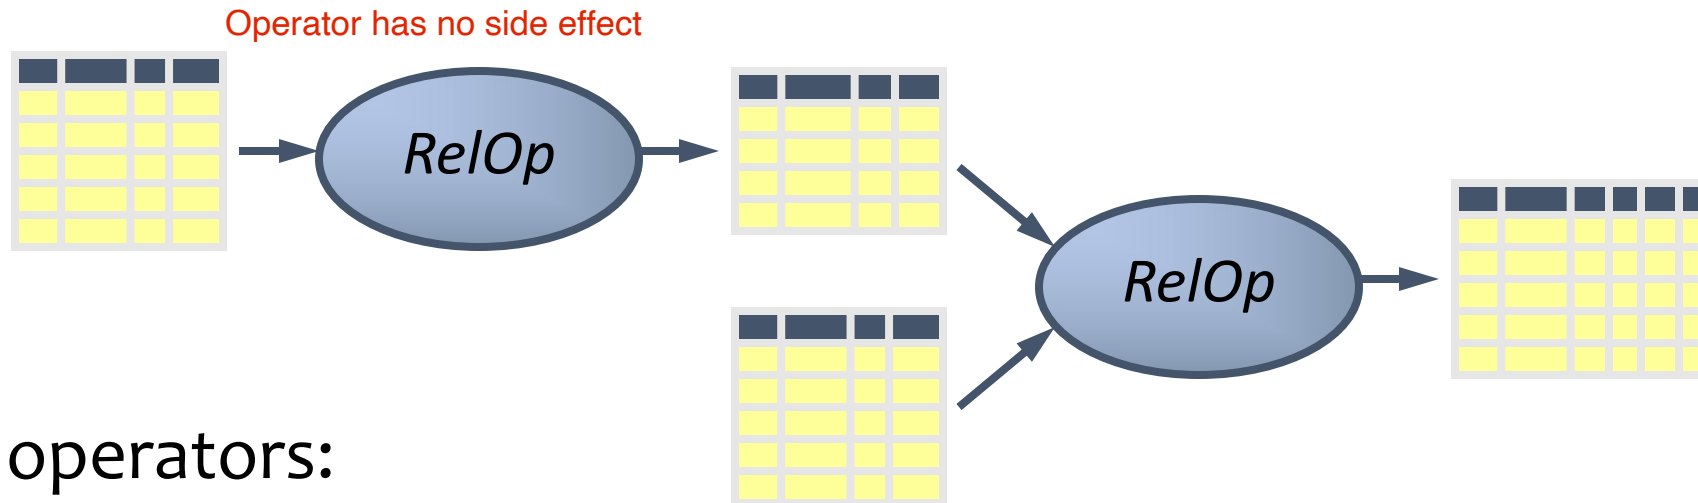
- Compare to **types** vs. collections of **objects of these types** in a programming language

# Example

- Schema *Ask the professor what he said regarding*
  - *User* (*uid* int, *name* string, *age* int, *pop* float)
  - *Group* (*gid* string, *name* string)
  - *Member* (*uid* int, *gid* string)
- Instance
  - *User*: {⟨142, Bart, 10, 0.9⟩, ⟨857, Milhouse, 10, 0.2⟩, ... }
  - *Group*: {⟨abc, Book Club⟩, ⟨gov, Student Government⟩, ... }
  - *Member*: {⟨142, dps⟩, ⟨123, gov⟩, ... }

# Relational algebra

A language for querying relational data based on “operators”



- **Core** operators:
  - Selection, projection, cross product, union, difference, and renaming
- Additional, **derived** operators:
  - Join, natural join, intersection, etc.
- Compose operators to make complex queries

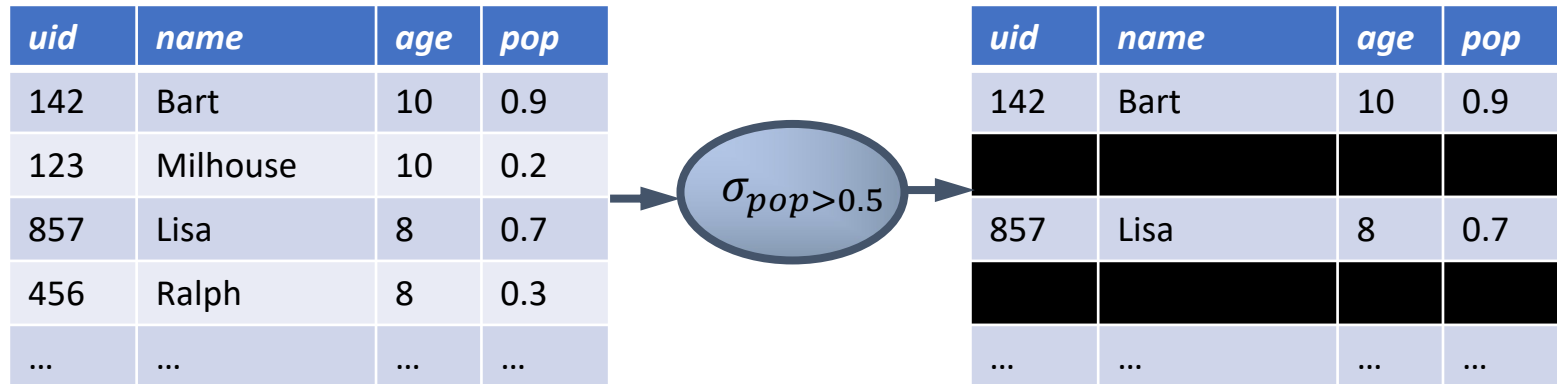
# Selection

- Input: a table  $R$
- Notation:  $\sigma_p R$ 
  - $p$  is called a **selection condition** (or **predicate**)
- Purpose: filter rows according to some criteria
- Output: same columns as  $R$ , but only rows **of**  $R$  that satisfy  $p$

# Selection example

- Users with popularity higher than 0.5

$$\sigma_{pop>0.5} User$$





# More on selection

- Selection condition can include any column of  $R$ , constants, comparison ( $=$ ,  $\leq$ , etc.) and Boolean connectives ( $\wedge$ : and,  $\vee$ : or,  $\neg$ : not)

- Example: users with popularity at least 0.9 and age under 10 or above 12

$$\sigma_{pop \geq 0.9 \wedge (age < 10 \vee age > 12)} User$$

- You must be able to evaluate the condition over **each single row** of the input table!

- Example: the most popular user

$$\sigma_{pop \geq \text{every pop in } User} User$$

**WRONG!**

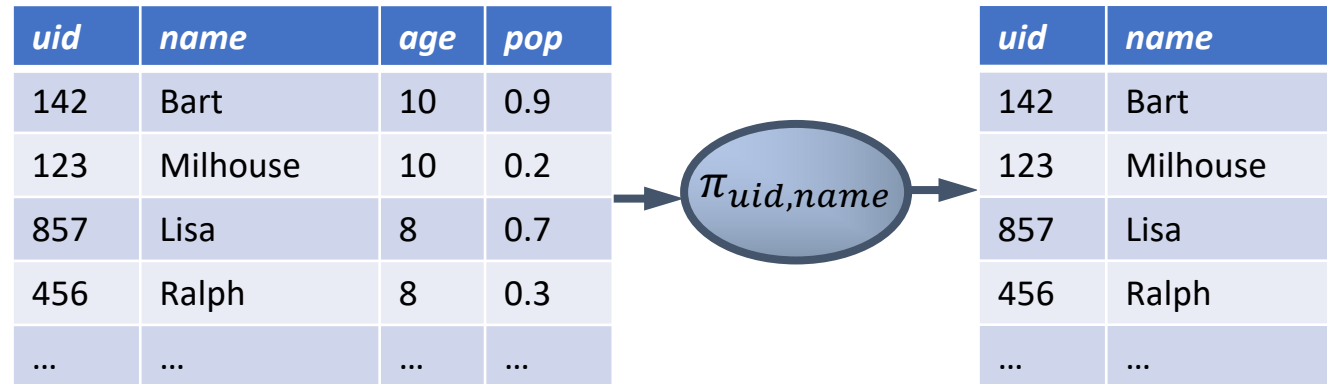
# Projection

- Input: a table  $R$
- Notation:  $\pi_L R$ 
  - $L$  is a list of columns in  $R$
- Purpose: output chosen columns
- Output: same rows, but only the columns in  $L$

# Projection example

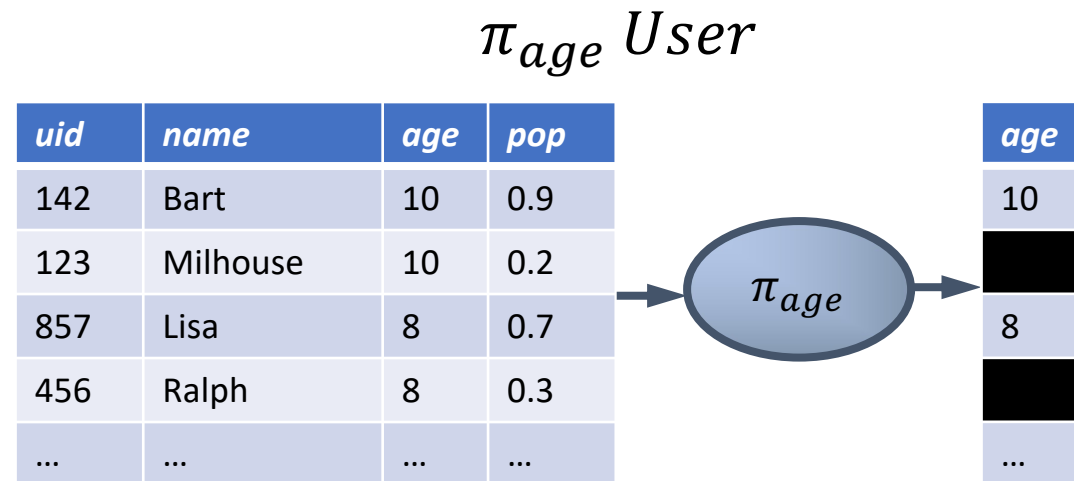
- IDs and names of all users

$\pi_{uid,name} User$



# More on projection

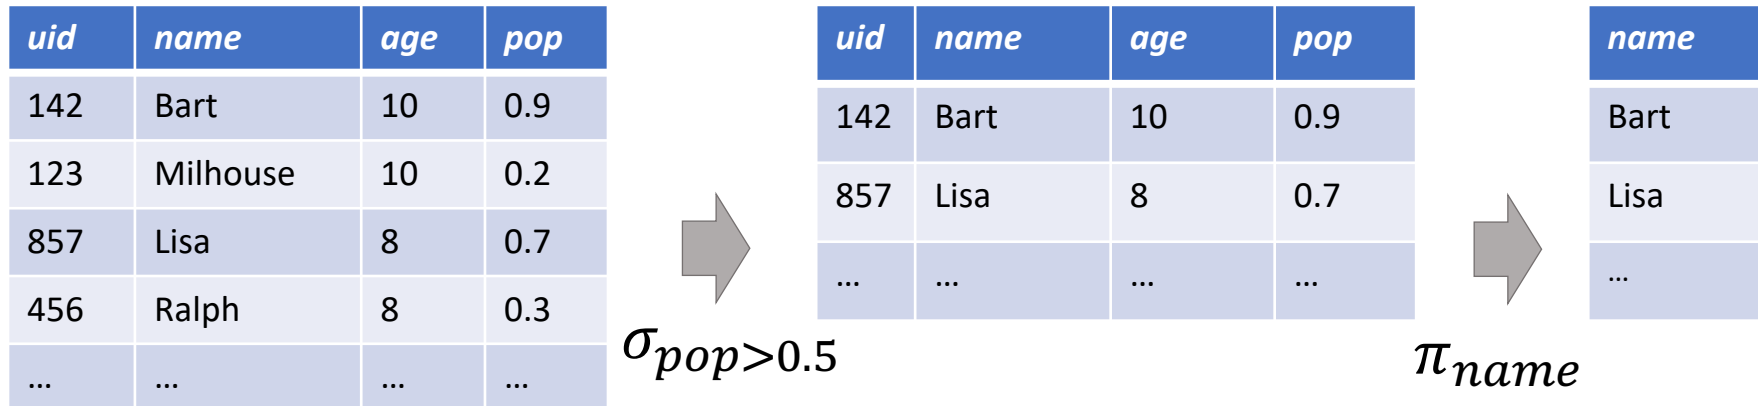
- Duplicate output rows are removed (by definition)
  - Example: user ages



# Composing Select and Project

- Names of users with popularity higher than 0.5

$$\pi_{name}(\sigma_{pop>0.5} User)$$

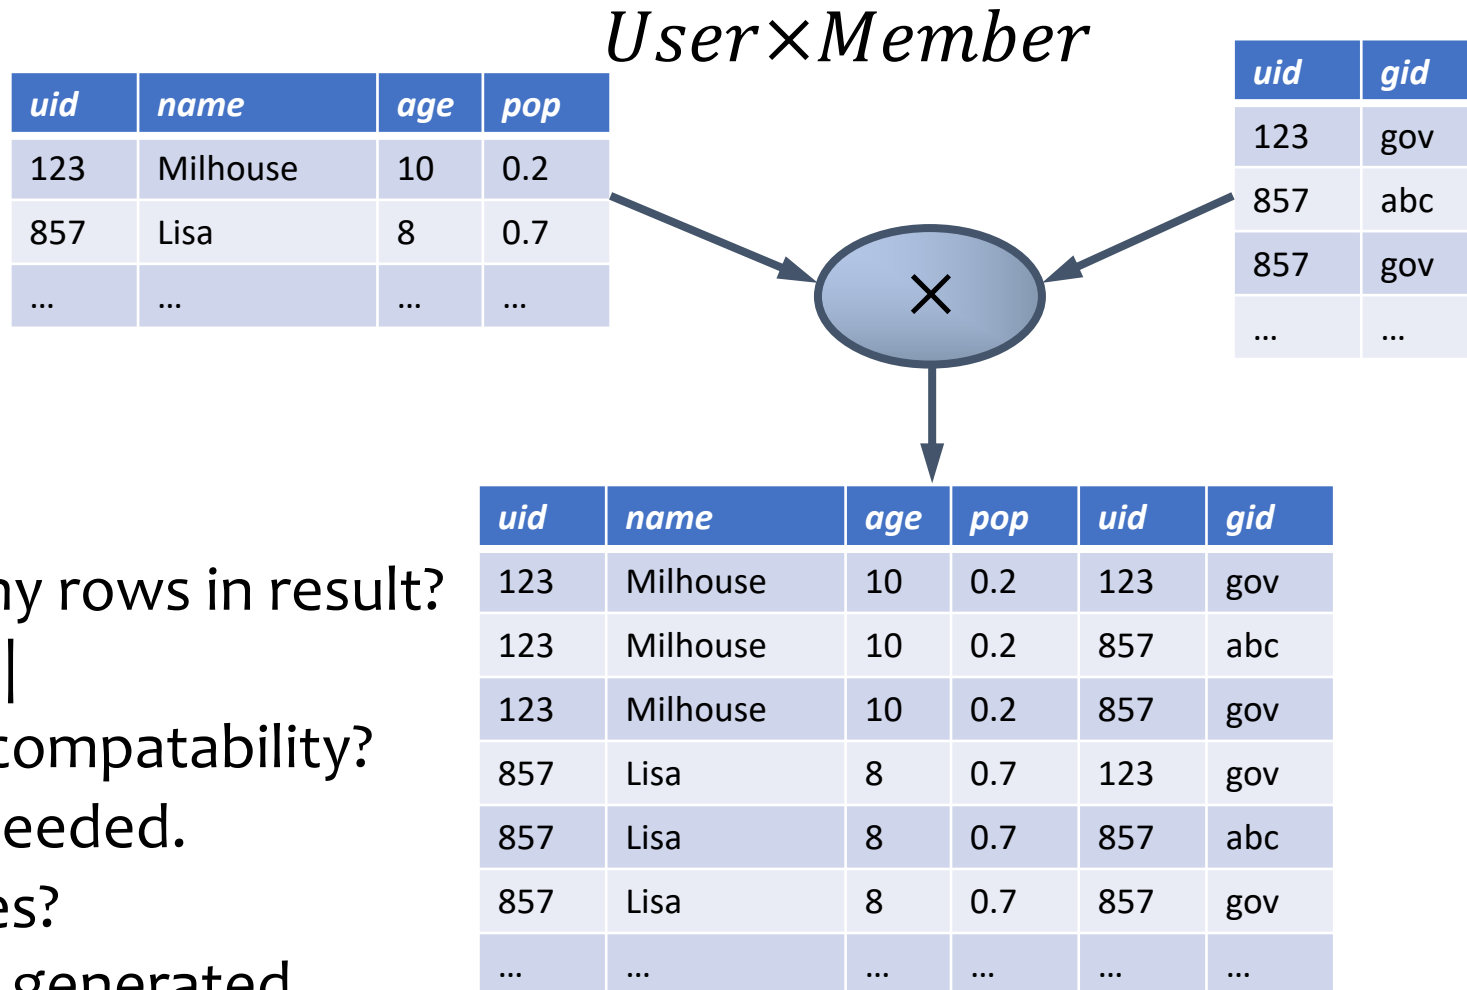


- What about:  $\sigma_{pop>0.5}(\pi_{name}(User))$ 
  - Invalid types. Input to  $\sigma_{pop>0.5}$  does not contain pop.

# Cross product

- Input: two tables  $R$  and  $S$
- Notation:  $R \times S$
- Purpose: pairs rows from two tables
- Output: for each row  $r$  in  $R$  and each  $s$  in  $S$ , output a row  $rs$  (concatenation of  $r$  and  $s$ )

# Cross product example



# A note on column ordering

- Ordering of columns is unimportant as far as contents are concerned

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>	<i>uid</i>	<i>gid</i>
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...	...	...	...	...	...

=

<i>uid</i>	<i>gid</i>	<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
123	gov	123	Milhouse	10	0.2
857	abc	123	Milhouse	10	0.2
857	gov	123	Milhouse	10	0.2
123	gov	857	Lisa	8	0.7
857	abc	857	Lisa	8	0.7
857	gov	857	Lisa	8	0.7
...	...	...	...	...	...

- So cross product is **commutative**, i.e., for any  $R$  and  $S$ ,  $R \times S = S \times R$  (up to the ordering of columns)



# Derived operator: join

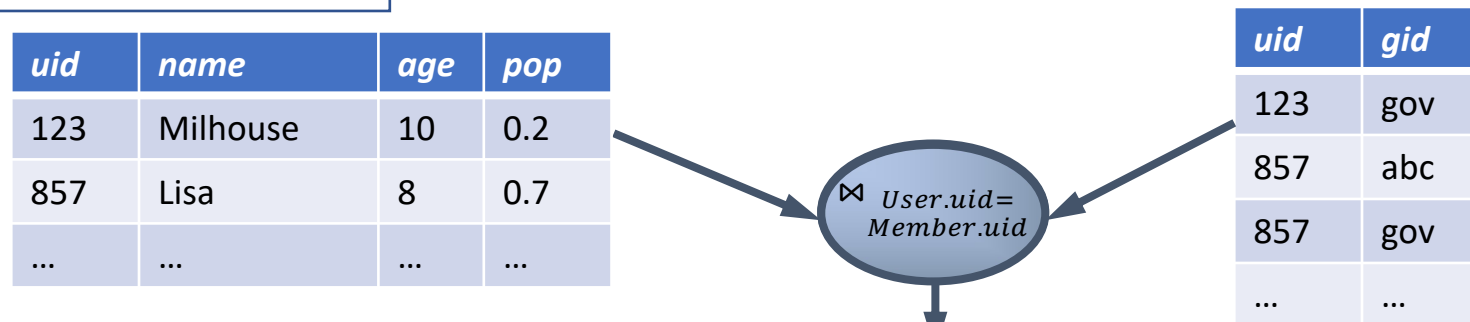
(A.k.a. “theta-join”)

- Input: two tables  $R$  and  $S$
- Notation:  $R \bowtie_p S$ 
  - $p$  is called a **join condition** (or **predicate**)
- Purpose: relate rows from two tables according to some criteria
- Output: for each row  $r$  in  $R$  and each row  $s$  in  $S$ , output a row  $rs$  if  $r$  and  $s$  satisfy  $p$
- Shorthand for  $\sigma_p(R \times S)$

# Join example

- Info about users, plus IDs of their groups

$User \bowtie_{User.uid=Member.uid} Member$



Prefix a column reference with table name and “.” to disambiguate identically named columns from different tables

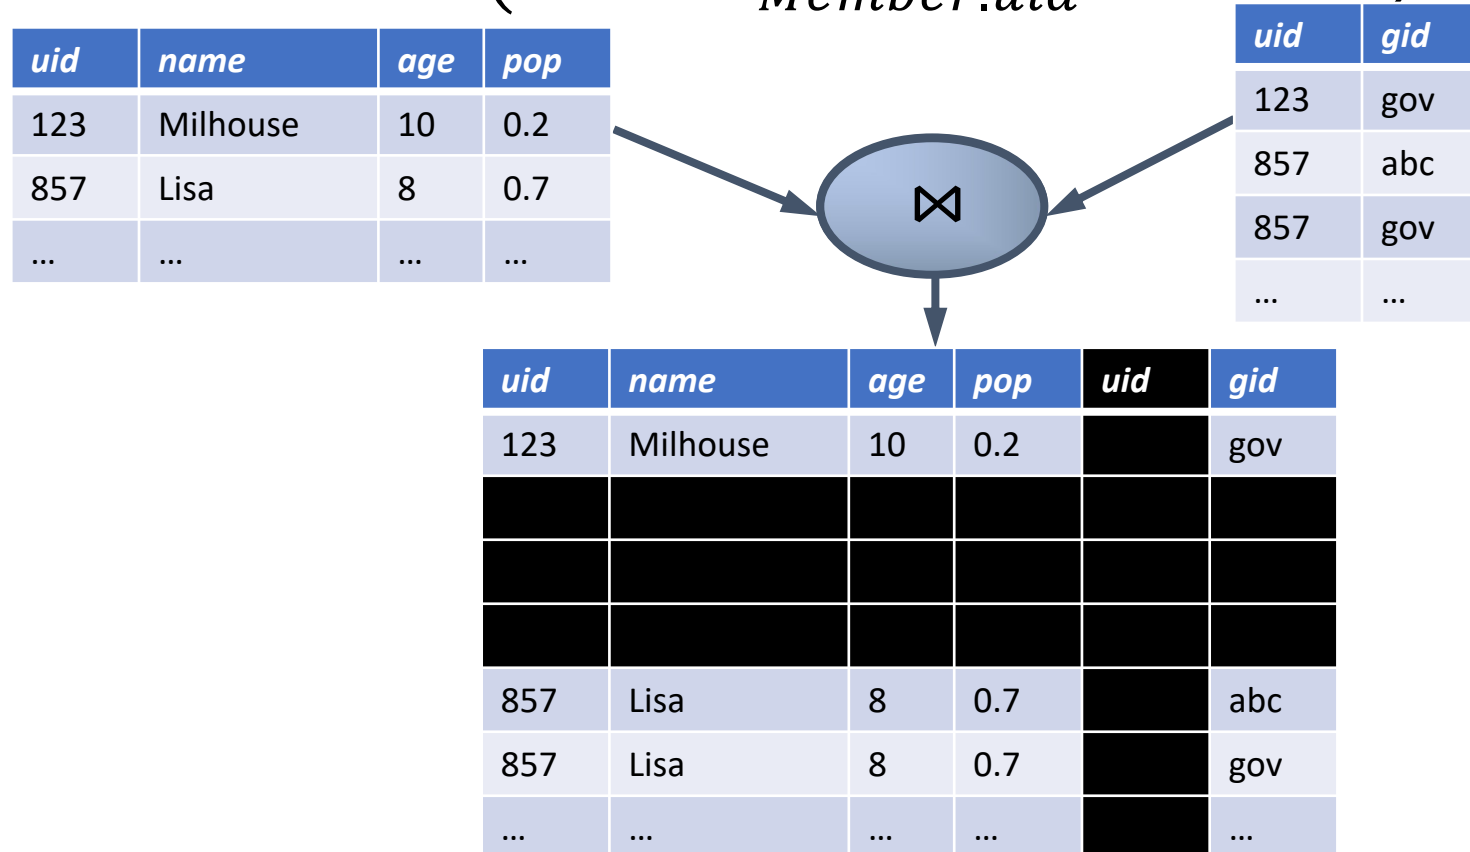
uid	name	age	pop	uid	gid
123	Milhouse	10	0.2	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...	...	...	...	...	...

# Derived operator: natural join

- Input: two tables  $R$  and  $S$
- Notation:  $R \bowtie S$
- Purpose: relate rows from two tables, and
  - Enforce equality between identically named columns
  - Eliminate one copy of identically named columns
- Shorthand for  $\pi_L(R \bowtie_p S)$ , where
  - $p$  equates each pair of columns common to  $R$  and  $S$
  - $L$  is the union of column names from  $R$  and  $S$  (with duplicate columns removed)

# Natural join example

$$\begin{aligned}
 User \bowtie Member &= \pi_{?}(User \bowtie_{?} Member) \\
 &= \pi_{uid, name, age, pop, gid} \left( User \bowtie_{\substack{User.uid = \\ Member.uid}} Member \right)
 \end{aligned}$$



# Union

- Input: two tables  $R$  and  $S$
- Notation:  $R \cup S$ 
  - $R$  and  $S$  must have identical schema
- Output:
  - Has the same schema as  $R$  and  $S$
  - Contains all rows in  $R$  and all rows in  $S$  (with duplicate rows removed)

<i>uid</i>	<i>gid</i>		<i>uid</i>	<i>gid</i>		<i>uid</i>	<i>gid</i>
123	gov	U	123	gov	=	123	gov
857	abc		857	gov		857	abc
						857	gov

# Difference

- Input: two tables  $R$  and  $S$
- Notation:  $R - S$ 
  - $R$  and  $S$  must have identical schema
- Output:
  - Has the same schema as  $R$  and  $S$
  - Contains all rows in  $R$  that are not in  $S$

<i>uid</i>	<i>gid</i>		<i>uid</i>	<i>gid</i>		<i>uid</i>	<i>gid</i>
123	gov	—	123	gov	=	857	abc
857	abc		857	gov			

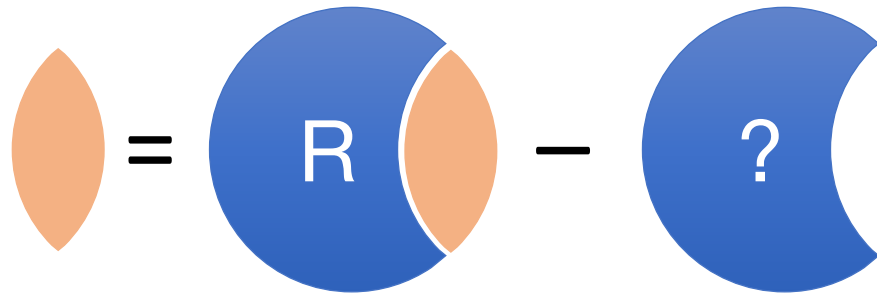

# Derived operator: intersection

- Input: two tables  $R$  and  $S$
- Notation:  $R \cap S$ 
  - $R$  and  $S$  must have identical schema
- Output:
  - Has the same schema as  $R$  and  $S$
  - Contains all rows that are in both  $R$  and  $S$
- Shorthand for? Could we also use the natural join?

<i>uid</i>	<i>gid</i>	$\cap$	<i>uid</i>	<i>gid</i>	$=$	<i>uid</i>	<i>gid</i>
123	gov		123	gov		123	gov
857	abc		857	gov			

# Intersection


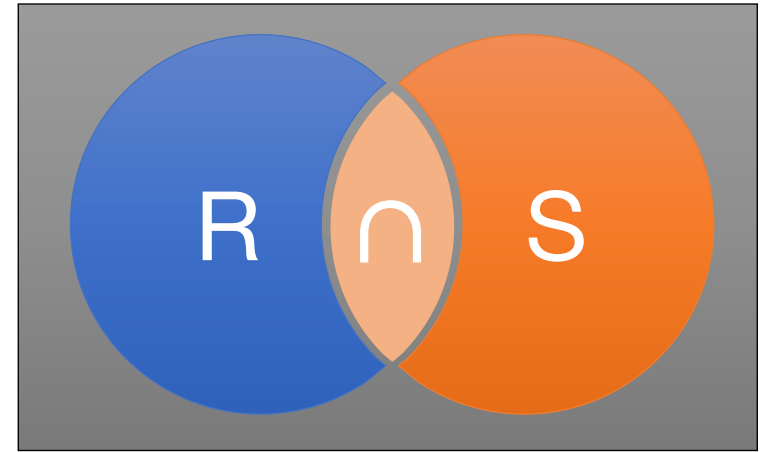
- $R \cap S = R - ?$





# Intersection

- $R \cap S = R - (R - S)$
- Also  $= S - (S - R)$



- How can you write it without –

$$R \bowtie S$$

# Renaming

- Input: a table  $R$
- Notation:  $\rho_S R$ ,  $\rho_{(A_1, A_2, \dots)} R$ , or  $\rho_{S(A_1, A_2, \dots)} R$
- Purpose: “rename” a table and/or its columns
- Output: a table with the same rows as  $R$ , but called differently

Renaming allows you to clarify which attributes belong to which instance

*Member*

<i>uid</i>	<i>gid</i>
123	gov
857	abc

$\rho_{M1(uid1, gid1)} Member =$

*M1*

<i>uid</i>	<i>gid</i>
123	gov
857	gov

Should the columns not be renamed?  
Professor: they should

# Renaming

- Used to
  - Avoid confusion caused by identical column names
  - Create identical column names for natural joins
- As with all other relational operators, it doesn't modify the database
  - Think of the renamed table as a copy of the original

*Member*

<i>uid</i>	<i>gid</i>
123	gov
857	abc

$$\rho_{M1(uid1,gid1)}Member =$$

*M1*

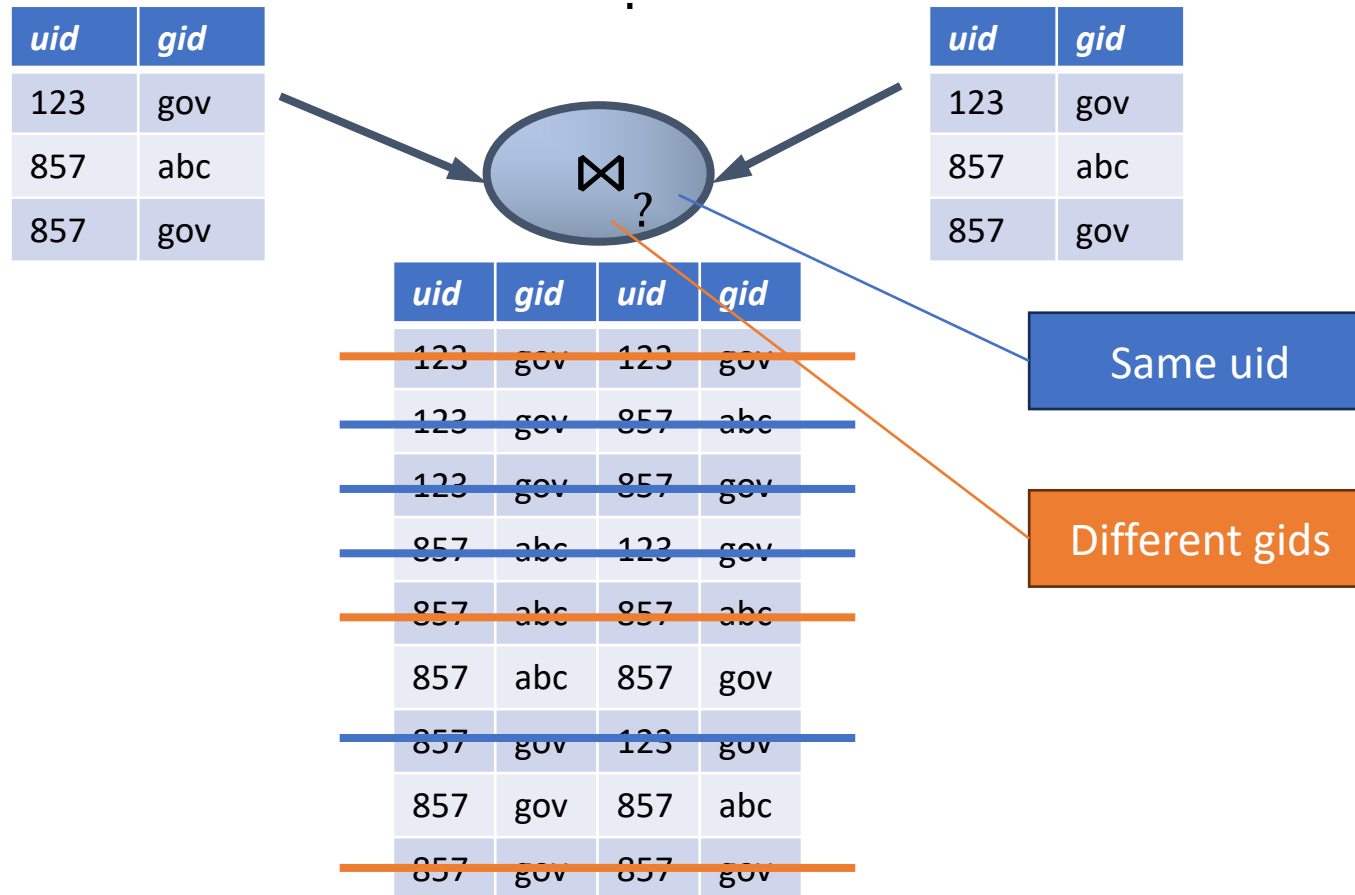
<i>uid</i>	<i>gid</i>
123	gov
857	gov

Note:

Learn to differentiate the natural join and regular join

# Renaming example

- IDs of users who belong to **at least two groups**  
*Member* ⋈<sub>?</sub> *Member*



# Renaming example

- IDs of users who belong to **at least two groups**

$Member \bowtie_? Member$

First condition will always be true

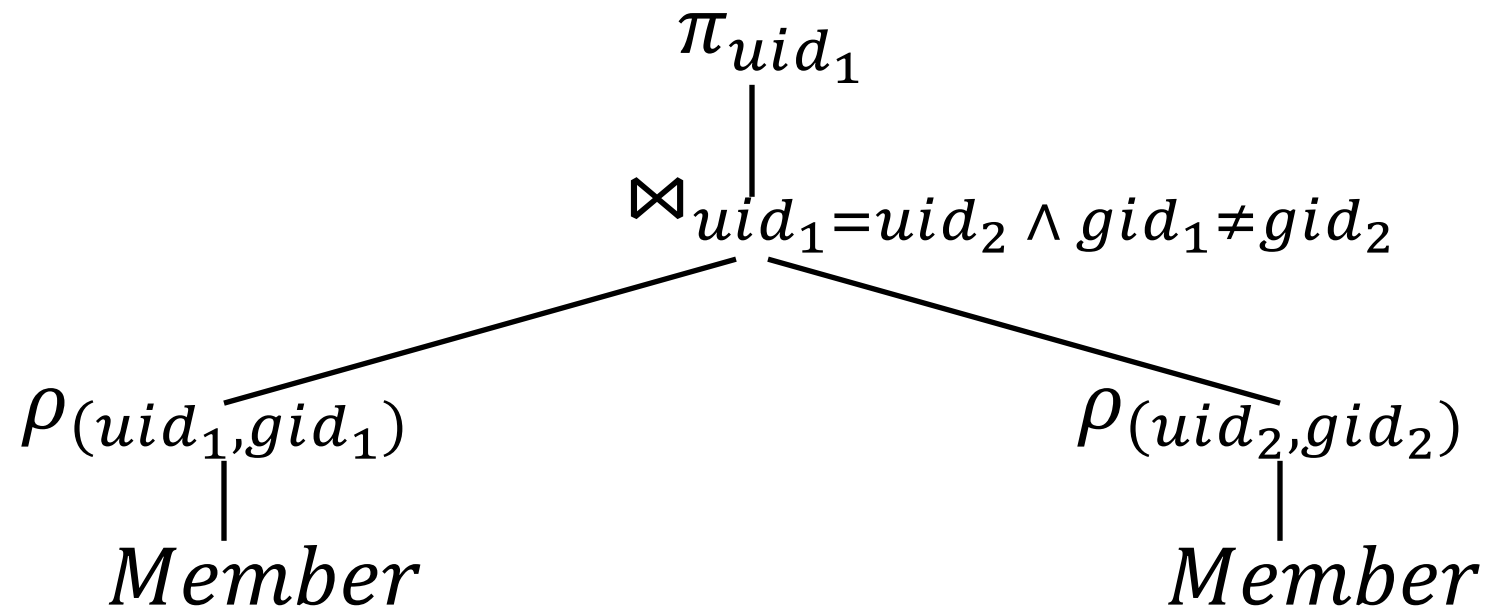
Second condition will always be false

$\pi_{uid} \left( Member \bowtie_{Member.uid=Member.uid \wedge Member.gid \neq Member.gid} Member \right)$

**WRONG!**

$$\pi_{uid_1} \left( \begin{array}{c} \rho_{(uid_1, gid_1)} Member \\ \bowtie_{uid_1=uid_2 \wedge gid_1 \neq gid_2} \\ \rho_{(uid_2, gid_2)} Member \end{array} \right)$$

# Expression tree notation



# Summary of core operators

- Selection:  $\sigma_p R$
- Projection:  $\pi_L R$
- Cross product:  $R \times S$
- Union:  $R \cup S$
- Difference:  $R - S$
- Renaming:  $\rho_{S(A_1, A_2, \dots)} R$ 
  - Does not really add “processing” power

# Summary of derived operators

- Join:  $R \bowtie_p S$
- Natural join:  $R \bowtie S$
- Intersection:  $R \cap S$
- Many more
  - Semijoin, anti-semijoin, quotient, ...

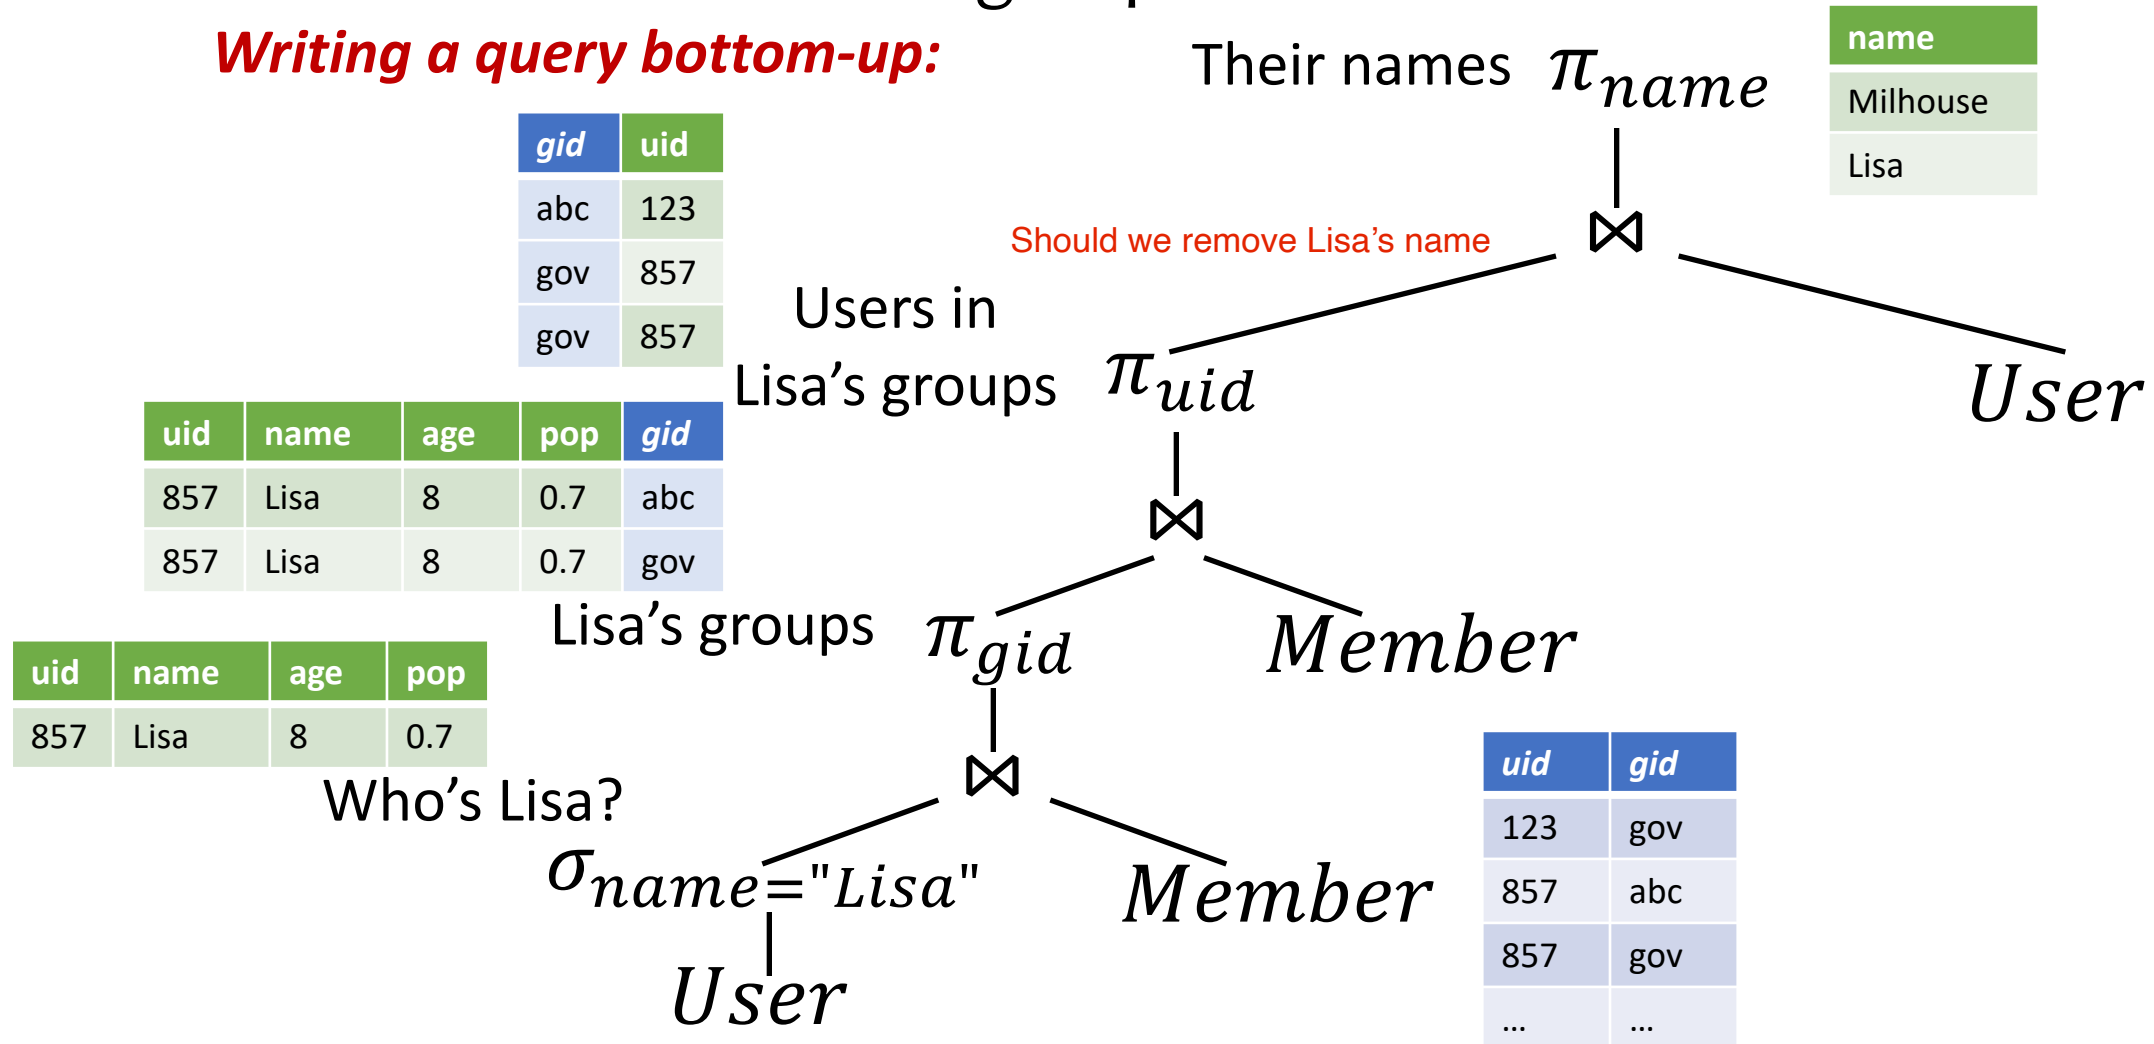


# An exercise

*User* (*uid* int, *name* string, *age* int, *pop* float)  
*Group* (*gid* string, *name* string)  
*Member* (*uid* int, *gid* string)

- Names of users in Lisa's groups

**Writing a query bottom-up:**

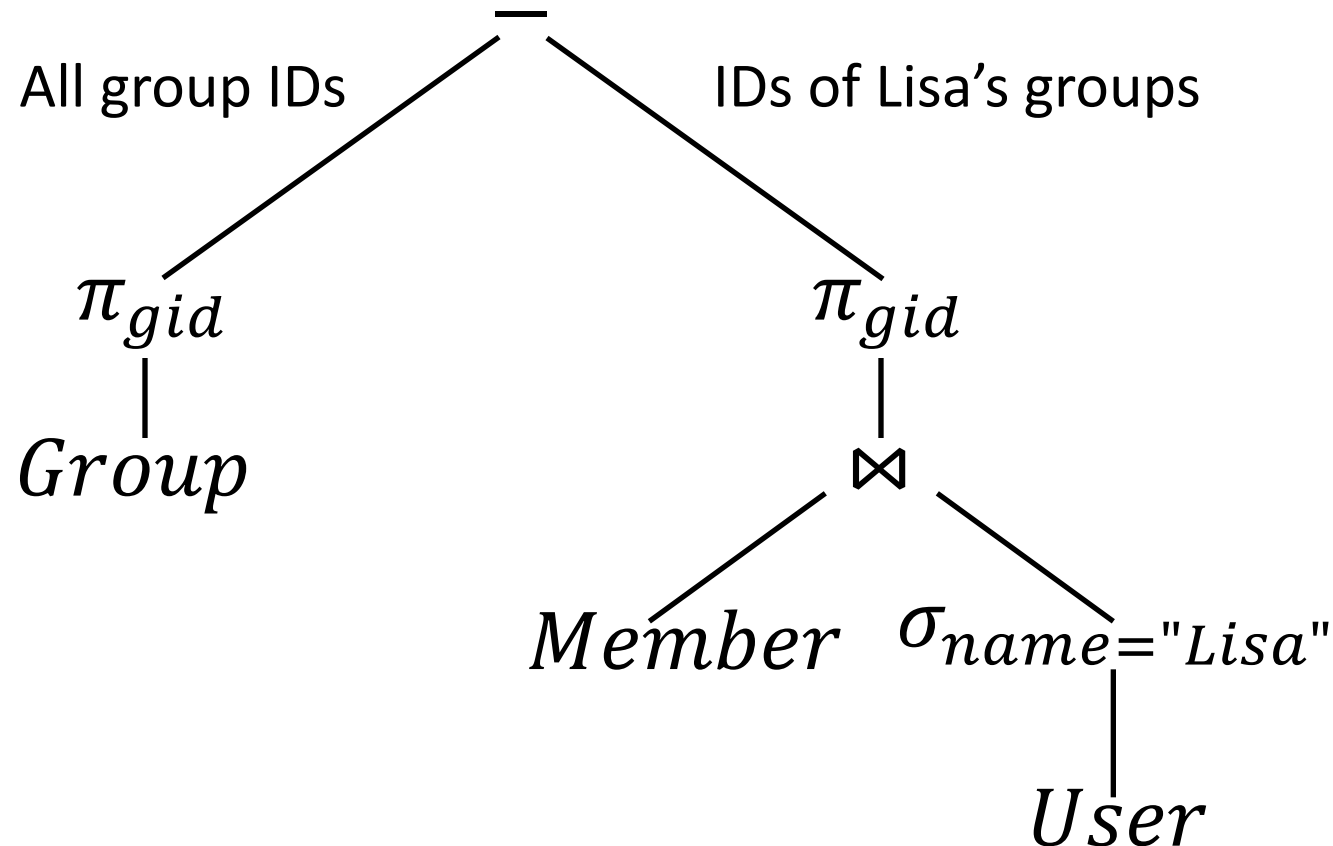


# Another exercise

*User* (uid int, name string, age int, pop float)  
*Group* (gid string, name string)  
*Member* (uid int, gid string)

- IDs of groups that Lisa doesn't belong to

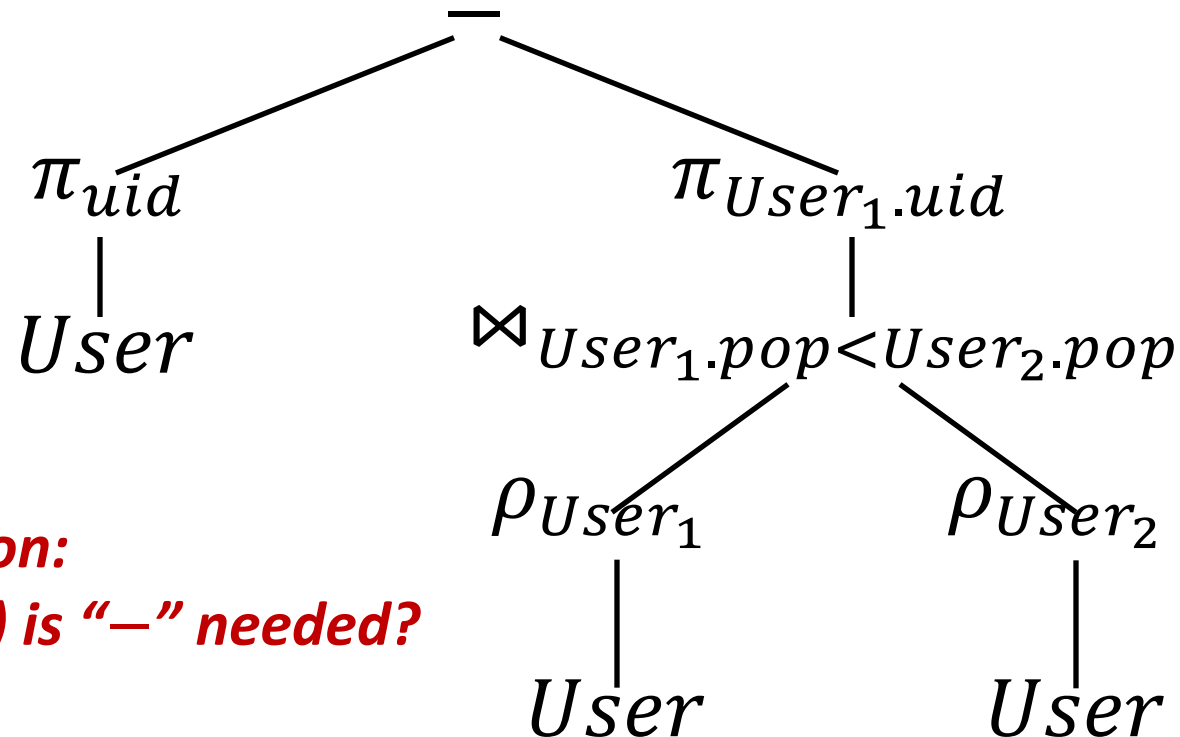
**Writing a query top-down:**



# A trickier exercise

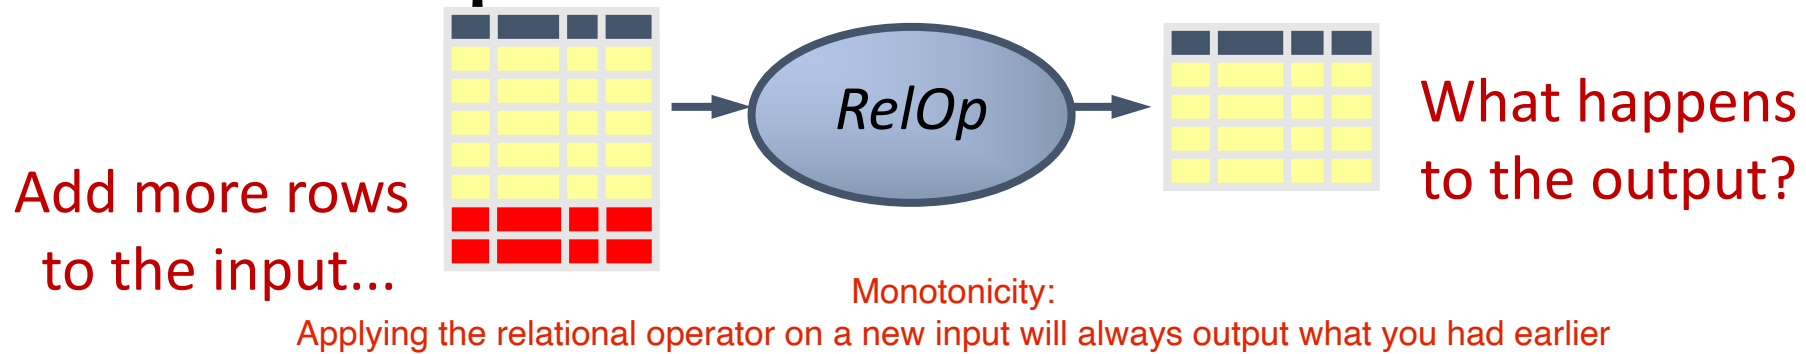
*User* (*uid* int, *name* string, *age* int, *pop* float)  
*Group* (*gid* string, *name* string)  
*Member* (*uid* int, *gid* string)

- Who are the most popular?
  - Who do NOT have the highest *pop* rating?
  - Whose *pop* is lower than somebody else's?



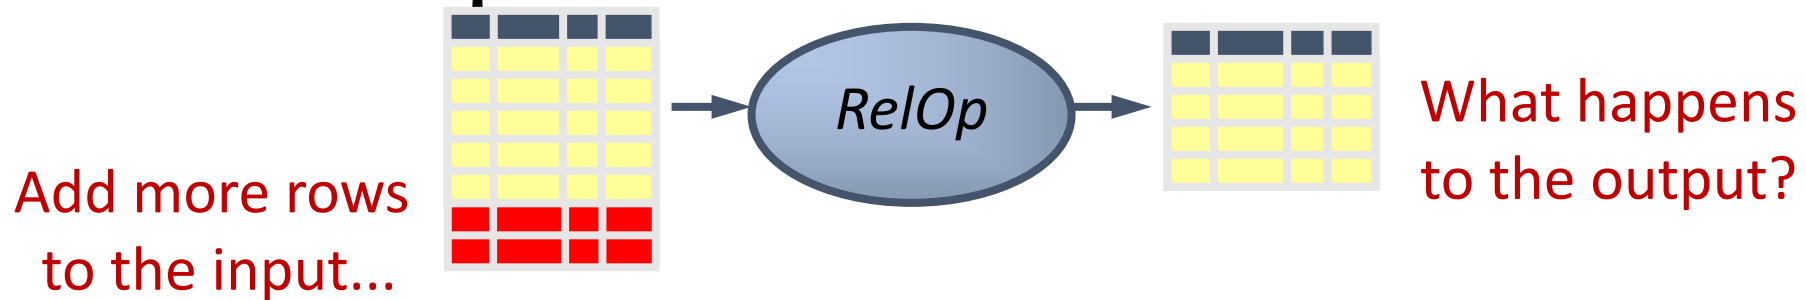
**A deeper question:**  
**When (and why) is “—” needed?**

# Monotone operators



- If some **old output rows** may need to be removed
    - Then the operator is **non-monotone**
  - Otherwise the operator is **monotone**
    - That is, old output rows always remain “correct” when more rows are added to the input
- Does monotonicity hold when removing rows from the input?  
Is there another property
- Formally, for a monotone operator  $op$ :  
 $R \subseteq R'$  implies  $op(R) \subseteq op(R')$  for any  $R, R'$

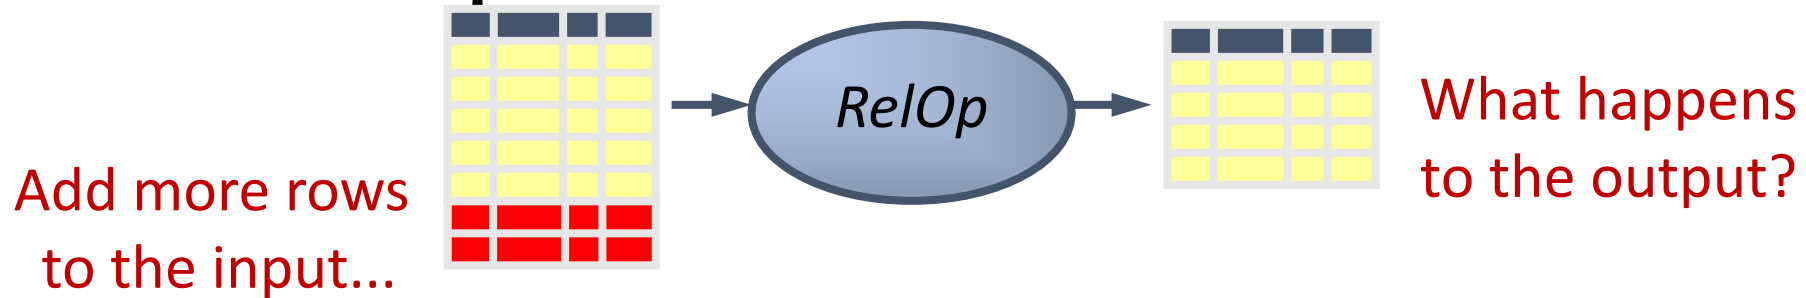
# Monotone operators



- If some **old output rows** may need to be removed
  - Then the operator is **non-monotone**
- Otherwise the operator is **monotone**
  - That is, old output rows always remain “correct” when more rows are added to the input

<i>uid</i>	<i>gid</i>	—	<i>uid</i>	<i>gid</i>	=	<i>uid</i>	<i>gid</i>
123	gov		123	gov		<del>857</del>	<del>abc</del>
857	abc		857	gov			
			857	abc			

# Monotone operators



- If some **old output rows** may need to be removed
  - Then the operator is **non-monotone**
- Otherwise the operator is **monotone**
  - That is, old output rows always remain “correct” when more rows are added to the input

<i>uid</i>	<i>gid</i>
123	gov
857	abc
857	gov
693	abc

—

<i>uid</i>	<i>gid</i>
123	gov
857	gov

=

<i>uid</i>	<i>gid</i>
857	abc
693	abc

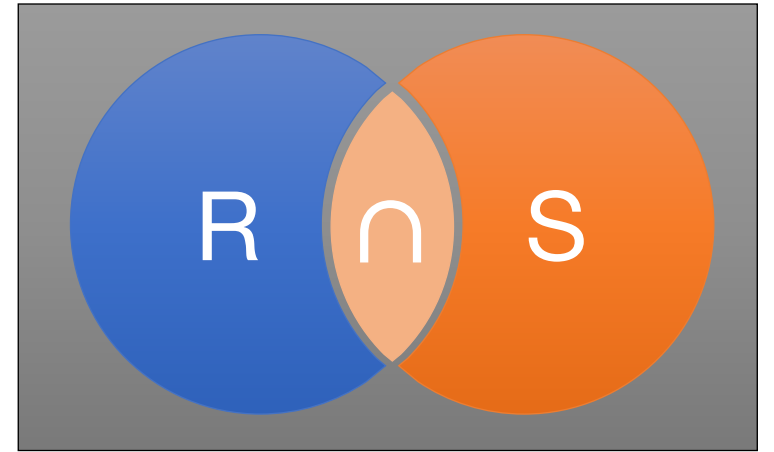
The old row is always “correct” no matter what is added to R

# Classification of relational operators

- Selection:  $\sigma_p R$  Monotone
- Projection:  $\pi_L R$  Monotone
- Cross product:  $R \times S$  Monotone
- Join:  $R \bowtie_p S$  Monotone
- Natural join:  $R \bowtie S$  Monotone
- Union:  $R \cup S$  Monotone
- Difference:  $R - S$  Monotone w.r.t.  $R$ ; non-monotone w.r.t  $S$
- Intersection:  $R \cap S$  Monotone?

# Is intersection monotonic?

- $R \cap S = R - (R - S)$



$$\text{Intersection} = R - (R - S)$$

The diagram illustrates the set identity  $R \cap S = R - (R - S)$  using Venn diagrams. On the left, a small orange lens-shaped region represents the intersection of two sets. This is followed by an equals sign. To the right of the equals sign, the expression is visualized: a blue circle labeled 'R' with a small orange segment on its right edge, minus (indicated by a minus sign) a large expression in parentheses. Inside the parentheses is a blue circle labeled 'R' with a yellow segment on its right edge, minus (indicated by a minus sign) an orange circle labeled 'S' with a yellow segment on its left edge. The yellow segments represent the parts of R and S that do not intersect.

- Yes!

$$R_1 \subseteq R_2 \Rightarrow S \cap R_1 \subseteq S \cap R_2$$



# Why is “–” needed for “highest”?

- Composition of monotone operators produces a **monotone query**
  - Old output rows remain “correct” when more rows are added to the input
- Is the “highest” query monotone?
  - No!
  - Current highest *pop* is 0.9
  - Add another row with *pop* 0.91
  - Old answer is invalidated
- So it must use difference!

# Why do we need core operator $X$ ?

- Difference
  - The only **non-monotone** operator
- Projection
  - The only operator that **removes columns**
- Cross product
  - The only operator that **adds columns**
- Union
  - The only operator that allows you to add rows?
  - A more rigorous argument?
- Selection?
  - Left as an exercise for the viewer

# Extensions to relational algebra

- Duplicate handling (“bag algebra”)
- Grouping and aggregation
- All these will come up when we talk about SQL
- But for now we will stick to standard relational algebra without these extensions

# Why is RA a good query language?

- Simple
  - A small set of core operators
  - Semantics are easy to grasp
- Declarative?
  - Yes, compared with older languages like CODASYL
  - Though assembling operators into a query does feel somewhat “procedural”
- Complete?
  - With respect to what?

# Relational Calculus

*User* (*uid* int, *name* string, *age* int, *pop* float)  
*Group* (*gid* string, *name* string)  
*Member* (*uid* int, *gid* string)

- $\{u.uid \mid u \in User \wedge \neg(\exists u' \in User: u.pop < u'.pop)\}$ , or
- $\{u.uid \mid u \in User \wedge (\forall u' \in User: u.pop \geq u'.pop)\}$
- Relational algebra = “safe” relational calculus
  - Every query expressible as a safe relational calculus query is also expressible as a relational algebra query
  - And vice versa
- Example of an “unsafe” relational calculus query
  - $\{u.name \mid \neg(u \in User)\}$
  - Cannot evaluate it just by looking at the database

# Relational Calculus

- RA and RC form the basis for “real” query languages (e.g. SQL), and for implementation
  - Relational Algebra
    - Queries are expressed by applying a sequence of operations to relations
    - More **operational/procedural**, very useful for representing how query executes
  - Relational Calculus
    - Let's users describe WHAT they want in first-order logic
    - Rather than HOW to compute it
    - Non-operational, **declarative**

# Codd's Theorem

- Established equivalence in expressivity between :
  - Relational Calculus
  - Relational Algebra
- Why an important result?
  - Connects declarative representation of queries with operational description
  - Constructive: we can “compile” SQL into relational algebra

# Limits of relational algebra

- Relational algebra has **no recursion**
  - Example: given relation *Friend*(*uid1*, *uid2*), who can Bart reach in his social network with any number of hops?
    - Writing this query in r.a. is impossible!
  - So RA is not as powerful as general-purpose languages (not Turing-complete)
- But why not?
  - Optimization becomes **undecidable**
  - Simplicity is empowering
  - Besides, you can always implement it at the application level (and recursion is added to SQL nevertheless 🤖)



# What's Next

- Next class:
  - Database design in E/R model (recording)
- Sign up Piazza
  - <https://piazza.com/sfu.ca/summer2024/cmpt354d100>
  - Access code: qc76kdzsz9l
- Assignment 1 releasing soon
- Stay tuned for RA help tools