

Database Systems I

CMPT 354 Summer 2024

Zhengjie Miao

Announcements (Fri. June 7)

- Midterm I (Wed. June 12)
 - Sample released
 - Just an example of the format and style
 - Exam will be 75 min (**not 80 min!**)
 - Covers Lec1-Lec8, including the part taught on Friday, June 7
 - Open-book, open-notes, but
 - **no collaboration**
 - **no laptop/phones/other electronics**
 - Please arrive early, the exam starts exactly at 3:35 and ends exactly at 4:50
 - Bring your SFU id

SQL features covered so far

- Query
 - SELECT-FROM-WHERE statements
 - Set and bag operations
 - Table expressions, subqueries
 - Aggregation and grouping
 - Ordering
 - NULL and Outerjoins
 - Modification & Constraints
- 👉 Today : triggers and views

Recap: Enforcing referential integrity

Example: *Member.uid* references *User.uid*

- Delete or update a *User* row whose *uid* is referenced by some *Member* row
 - Option 1: Reject
 - Option 2: Cascade --- ripple changes to all referring rows

<i>uid</i>	<i>name</i>	...
142	Bart	...
123	Milhouse	...
857	Lisa	...
456	Ralph	...
789	Nelson	...
...

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

```
CREATE TABLE Member
(uid INT NOT NULL REFERENCES
User(uid) ON DELETE CASCADE,
...);
```

“Active” data

- Constraint enforcement: When an operation violates a constraint, abort the operation or try to “fix” data
 - Example: enforcing referential integrity constraints
 - Generalize to arbitrary constraints?
- Data monitoring: When something happens to the data, automatically execute some action
 - Example: When price rises above \$20 per share, sell
 - Example: When enrollment is at the limit and more students try to register, email the instructor

Triggers

- A **trigger** is an event-condition-action (ECA) rule
 - When **event** occurs, test **condition**; if condition is satisfied, execute **action**
 - Different DBMS support different syntax, but concepts remain the same
 - E.g., PostgreSQL syntax \neq what we'll present here
- Example:
 - **Event**: some user's popularity is updated
 - **Condition**: the user is a member of cks ("Cool Kids") and *pop* drops below 0.5
 - **Action**: kick that user out of cks



Triggers

- A **trigger** is an event-condition-action (ECA) rule
 - When **event** occurs, test **condition**; if condition is satisfied, execute **action**
 - Different DBMS support different syntax, but concepts remain the same
 - E.g., PostgreSQL syntax \neq what we'll present here

Delete/update a row in User

Whether its uid is
referenced by some row in
Member

If Yes: reject/delete/
cascade/NULL

Referential constraints

Event



Condition



Action

Some user's popularity is updated

The user is a member of cks ("Cool Kids") and *pop* drops below 0.5

If Yes: kick that user out of cks

Data Monitoring

Trigger example

Review the syntax for a trigger

```
CREATE TRIGGER PickyCKS Event  
AFTER UPDATE OF pop ON User  
REFERENCING NEW ROW AS newUser  
FOR EACH ROW
```

```
WHEN (newUser.pop < 0.5)  
    AND (newUser.uid IN (SELECT uid  
                        FROM Member  
                        WHERE gid = 'cks')) Condition
```

```
DELETE FROM Member  
WHERE uid = newUser.uid AND gid = 'cks'; Action
```


Trigger option 1

- Possible events include:
 - **INSERT ON** table, **DELETE ON** table, **UPDATE [OF column] ON** table

```
CREATE TRIGGER PickyCKS
AFTER UPDATE OF pop ON User
REFERENCING NEW ROW AS newUser
FOR EACH ROW
```

Event

```
WHEN (newUser.pop < 0.5)
    AND (newUser.uid IN (SELECT uid
                        FROM Member
                        WHERE gid = 'cks'))
```

Condition

```
DELETE FROM Member
WHERE uid = newUser.uid AND gid = 'cks';
```

Action

Trigger option 2

- Timing—action can be executed:
 - **AFTER** or **BEFORE** the triggering event
 - **INSTEAD OF** the triggering event on views (more later)

```
CREATE TRIGGER PickyCKS Event  
AFTER UPDATE OF pop ON User  
REFERENCING NEW ROW AS newUser  
FOR EACH ROW
```

```
WHEN (newUser.pop < 0.5)  
    AND (newUser.uid IN (SELECT uid  
                        FROM Member  
                        WHERE gid = 'cks')) Condition
```

```
DELETE FROM Member  
WHERE uid = newUser.uid AND gid = 'cks'; Action
```

Trigger option 2

- Timing—action can be executed:
 - **AFTER** or **BEFORE** the triggering event
 - **INSTEAD OF** the triggering event on views (more later)

```
CREATE TRIGGER NoFountainOfYouth
BEFORE UPDATE OF age ON User
REFERENCING OLD ROW AS o,
              NEW ROW AS n
FOR EACH ROW
WHEN (n.age < o.age)
SET n.age = o.age;
```

Event

Condition

Action

- BEFORE triggers are often used to “condition” data
- Another option is to raise an error in the trigger body to abort the transaction that caused the trigger to fire

Trigger option 3 (Row-level)

- Granularity—trigger can be activated:
 - **FOR EACH ROW** modified

```
CREATE TRIGGER PickyCKS
AFTER UPDATE OF pop ON User
REFERENCING NEW ROW AS newUser
FOR EACH ROW
```

Event

```
WHEN (newUser.pop < 0.5)
AND (newUser.uid IN (SELECT uid
                     FROM Member
                     WHERE gid = 'cks'))
```

Condition

```
DELETE FROM Member
WHERE uid = newUser.uid AND gid = 'cks';
```

Action

Trigger option 3 (Statement-level)

- Granularity—trigger can be activated:
 - **FOR EACH ROW** modified
 - **FOR EACH STATEMENT** that performs modification

```
CREATE TRIGGER PickyCKS
AFTER UPDATE OF pop ON User
REFERENCING NEW TABLE AS newUsers
FOR EACH STATEMENT
```

Event

```
DELETE FROM Member
WHERE gid = 'cks'
  AND (uid IN (SELECT uid
                FROM newUsers
                WHERE pop < 0.5));
```

Condition & Action

Transition variables

- **OLD ROW**: the modified row before the triggering event
- **NEW ROW**: the modified row after the triggering event
- **OLD TABLE**: a hypothetical read-only table containing all rows to be modified before the triggering event
- **NEW TABLE**: a hypothetical table containing all modified rows after the triggering event
- Not all of them make sense all the time, e.g.
 - BEFORE/AFTER INSERT statement-level triggers
 - Can use only NEW TABLE
 - BEFORE/AFTER UPDATE row-level triggers
 - Can use only OLD ROW and NEW ROW
 - BEFORE/AFTER DELETE row-level triggers
 - Can use only OLD ROW
 - etc.

Why hypothetical tables?
Is the new table hypothetical since we
may reject it?

Statement- vs. row-level triggers

Review statement and row-level triggers

Why are both needed?

- Certain triggers are only possible at statement level
 - If the number of users inserted by this statement exceeds 100 and their average age is below 13, then ...
- Simple row-level triggers are easier to implement
 - Statement-level triggers require significant amount of state to be maintained in OLD TABLE and NEW TABLE
 - However, a row-level trigger gets fired for each row, so complex row-level triggers may be less efficient for statements that modify many rows

System issues

- Recursive firing of triggers
 - Action of one trigger causes another trigger to fire
 - Can get into an infinite loop
- Interaction with constraints (tricky to get right!)
 - When do we check if a triggering event violates constraints?
 - After a BEFORE trigger (so the trigger can fix a potential violation)
 - Before an AFTER trigger
 - AFTER triggers also see the effects of, say, cascaded deletes caused by referential integrity constraint violations

Views

- A **view** is like a “virtual” table
 - Defined by a query, which describes how to compute the view contents on the fly
 - DBMS stores the **view definition query** instead of view contents
 - Can be used in queries just like a regular table
- In contrast to temporary tables defined by WITH (visible only in the same statement), views are **part of the schema** and visible to all statements

Creating and dropping views

- Example: members of Cool Kids
 - `CREATE VIEW CoolKids AS
SELECT * FROM User
WHERE uid IN (SELECT uid FROM Member
WHERE gid = 'cks');`
 - Tables used in defining a view are called “base tables”
 - E.g., *User* and *Member* above
- To drop a view
 - `DROP VIEW CoolKids;`

Using views in queries

- Example: find the average popularity of members in Cool Kids

- `SELECT AVG(pop) FROM CoolKids;`

- To process the query, replace the reference to the view by its definition

- `SELECT AVG(pop)
FROM (SELECT * FROM User
 WHERE uid IN
 (SELECT uid FROM Member
 WHERE gid = 'cks'))
AS CoolKids;`

Why use views?

- To hide data from users
 - To hide complexity from users
 - **Logical data independence**
 - If applications deal with views, we can change the underlying schema without affecting applications
 - Recall **physical data independence**: change the physical organization of data without affecting applications
 - To provide a uniform interface for different implementations or sources
- ☞ Real database applications use tons of views

Modifying views

- Does it even make sense, since views are virtual?
- It does make sense if we want users to really see views as tables
- Goal: modify the base tables such that the modification would appear to have been accomplished on the view

A simple case

```
CREATE VIEW UserPop AS  
  SELECT uid, pop FROM User;
```

```
DELETE FROM UserPop WHERE uid = 123;
```

translates to:

```
DELETE FROM User WHERE uid = 123;
```

An impossible case

```
CREATE VIEW PopularUser AS  
  SELECT uid, pop FROM User  
  WHERE pop >= 0.8;
```

```
INSERT INTO PopularUser  
  VALUES(987, 0.3);
```

- No matter what we do on *User*, the inserted row will not be in *PopularUser*

A case with too many possibilities

```
CREATE VIEW AveragePop(pop) AS  
SELECT AVG(pop) FROM User;
```

- Note that you can rename columns in view definition

```
UPDATE AveragePop SET pop = 0.5;
```

- Set everybody's *pop* to 0.5?
- Adjust everybody's *pop* by the same amount?
- Just lower the Cool Kid's *pop*?

SQL92 updateable views

- More or less just single-table selection queries
 - No join
 - No aggregation
 - No subqueries
- Arguably somewhat restrictive
- Still might get it wrong in some cases
 - See the slide titled “An impossible case”
 - Adding **WITH CHECK OPTION** to the end of the view definition will make DBMS reject such modifications

INSTEAD OF triggers for views

```
CREATE VIEW AveragePop(pop) AS  
SELECT AVG(pop) FROM User;
```

```
CREATE TRIGGER AdjustAveragePop  
INSTEAD OF UPDATE ON AveragePop  
REFERENCING OLD ROW AS o,  
              NEW ROW AS n  
FOR EACH ROW  
UPDATE User  
SET pop = pop + (n.pop-o.pop);
```

- What does this trigger do?

...	<i>pop</i>
	0.9
	0.7
	0.3
	0.5


```
UPDATE AveragePop SET pop = 0.5;
```

INSTEAD OF triggers for views

```
CREATE VIEW AveragePop(pop) AS  
SELECT AVG(pop) FROM User;
```

```
CREATE TRIGGER AdjustAveragePop  
INSTEAD OF UPDATE ON AveragePop  
REFERENCING OLD ROW AS o,  $\xrightarrow{0.6}$   
NEW ROW AS n  $\xrightarrow{0.5}$   
FOR EACH ROW  
UPDATE User  
SET pop = pop + (n.pop - o.pop);
```

- What does this trigger do?



```
UPDATE AveragePop SET pop = 0.5;
```

SQL features covered so far

- Query
- Modification
- Constraints
- Triggers
- Views