

Database Systems I

CMPT 354 Summer 2024

Zhengjie Miao

Announcements (Fri. June 28)


- Assignment 3 due tonight
- Assignment 4 released yesterday
 - Due July 8
 - You may refer to [the tips](#) and install MongoDB on your machine
 - Online tester is also available
 - May not as stable as for RA/SQL queries
 - https://ratest.cs.sfu.ca/mongo_test

JSON (JavaScript Object Notation)

- Very lightweight data exchange format
 - Much less verbose and easier to parse than XML
 - Increasingly used for data exchange over Web: many Web APIs use JSON to return responses/results
- Based on JavaScript
 - Conforms to JavaScript object/array syntax—you can directly manipulate JSON representations in JavaScript
- But it has gained widespread support by all programming languages

Complaints about XML format:

Repetition of elements and needing to put the opening/closing braces



Example JSON vs. XML

```
[
  {
    "ISBN": "ISBN-10",
    "price": 80.00,
    "title": "Foundations of Databases",
    "authors": [ "Abiteboul", "Hull", "Vianu" ],
    "publisher": "Addison Wesley",
    "year": 1995,
    "sections": [
      {
        "title": "Section 1",
        "sections": [
          { "title": "Section 1.1" },
          { "title": "Section 1.2" }
        ]
      },
      { "title": "Section 2" }
    ]
  },
  ...
]
```

```
<bibliography>
  <book ISBN="ISBN-10" price="80.00">
    <title>Foundations of Databases</title>
    <author>Abiteboul</author>
    <author>Hull</author>
    <author>Vianu</author>
    <publisher>Addison Wesley</publisher>
    <year>1995</year>
    <section>
      <title>Section 1</title>
      <section><title>Section 1.1</title></section>
      <section><title>Section 1.2</title></section>
    </section>
    <section>
      <title>Section 2</title>
    </section>
  </book>
</bibliography>
```

JSON data model

```
[
  { "ISBN": "ISBN-10",
    "price": 80.00,
    "title": "Foundations of Databases",
    "authors": [ "Abiteboul", "Hull", "Vianu" ],
    "publisher": "Addison Wesley",
    "year": 1995,
    "sections": [
      { "title": "Section 1",
        "sections": [
          { "title": "Section 1.1" },
          { "title": "Section 1.2" }
        ]
      },
      { "title": "Section 2" }
    ]
  }, ...
]
```

- Two basic constructs
 - **Array**: comma-separated list of “things” enclosed by brackets
 - Order is important
 - **Object**: comma-separated set of pairs enclosed by braces; each pair consists of an attribute name (string) and a value (any “thing”)
 - Order is unimportant
 - Attribute names “should” be unique within an object
- Simple types: numbers, strings (in double quotes), and special values **true**, **false**, and **null**
- Thing = a **simple value** or an **array** or an **object**

JSON Schema

- Recall the advantages of having a schema
 - Defines a structure, helps catch errors, facilitates exchange/automation, informs optimization...
- Just like relational data and XML, JSON is getting a schema standard too!
 - Up and coming, but still a draft at this stage

```
{  
  "definitions": {  
    "sections": {  
      "type": "array",  
      "description": "Sections.",  
      "sections": {"$ref": "#definitions/sections"},  
      "minItems": 0  
    }  
  },  
  "title": "Book",  
  "type": "object",  
  "properties": {  
    "ISBN": {  
      "type": "string",  
      "description": "The book's ISBN number."  
    },  
    "price": {  
      "type": "number",  
      "description": "The book's price.",  
      "exclusiveMinimum": 0  
    },  
    "...": {  
      "$ref": "#definitions/sections",  
    }  
  }  
}
```

MongoDB

Around 2000, people wanted to have more concurrency,
so the level of isolation had to be lowered

- One of the “NoSQL” poster children
- Started in 2007
- Targeting semi-structured data in JSON
- Designed to be easy to “scale out”
- Good support for indexing, partitioning, replication
- Nice integration in Web development stacks
- Not-so-great support for joins (or complex queries) or transactions



Inside a MongoDB database

- **Database** = a number of “collections”
- **Collection** = a list of “documents”
- **Document** = a JSON object
 - Must have an `_id` attribute whose value can uniquely identify a document within the collection
- ☞ In other words, a database has collections of similarly structured “documents”
 - Much like tables of records, as opposed to one big XML document that contains all data
- Good reads
 - <https://www.mongodb.com/resources/languages/json-schema-examples>

Querying MongoDB

- `find()` and `sort()`
 - Analogous to single-table selection/projection/sort
- “Aggregation” pipeline
 - With “stages” analogous to relational operators
 - Join, group-by, restructuring, etc.
- MapReduce (now deprecated):
 - Supports user-defined functions
- We won’t cover syntax for creating/updating MongoDB databases in lecture
 - Read the tips we provide and the manuals!

Key features to look out for

- Queries written as JSON objects themselves!
 - Natural in some cases (e.g., for specifying conditions on subsets of attributes), but awkward/misleading in others
- Simple path expressions using the “dot notation”
 - Analogous to XPath “/”
- Arrays within objects
 - Work on nested array directly using constructs like dot-index notation, `$elemMatch`, `$map`, and `$filter`
 - Or “**unnest**” an array so its elements get paired with the owner object in turn for pipeline processing
 - A fundamental concept in working with nested data

Basic MongoDB find()

- All books
`db.bib.find()`
- Books with title “Foundations of Databases”
`db.bib.find({ title: "Foundations of Databases" })`
- Books whose title contains “Database” or “database” and whose price is lower than \$50
`db.bib.find({ title: /[dD]atabase/, price: {$lt:50} })`

- Assume db refers to the database and db.bib refers to the collection of books
- In A4 you will need a couple more calls to pretty-print the output

Basic MongoDB find()

- All books
`db.bib.find()`
- Books with title “Foundations of Databases”
`db.bib.find({ title: "Foundations of Databases" })`
- Books whose title contains “Database” or “database” and whose price is lower than \$50
`db.bib.find({ title: /[dD]atabase/, price: {$lt:50} })`
- Books with price between \$70 and \$100
`db.bib.find({$and: [{price: {$gte:70}}, {price: {$lte:100}}]})`
 - By the way, why wouldn't the following work?
`db.bib.find({ price: {$gte:70}, price: {$lte:100} })`
- Books authored by Widom
`db.bib.find({ authors: "Widom" })`
 - Note the implicit existential quantification

- Assume db refers to the database and db.bib refers to the collection of books
- In A4 you will need a couple more calls to pretty-print the output

The second value for price will overwrite the first value, obtaining books less than or equal to 100 dollars

No general “twig” matching!

- Suppose for a moment publisher is an object itself, with attributes name, state, and country
- The following query won't get you database books by US publishers:

```
db.bib.find({ title: /[dD]atabase/,  
             publisher: { country: "US" } })
```

- Instead, the condition on publisher is satisfied only if it is an object with exactly one attribute, and this attribute must be named country and has value "US"
- What happens is that MongoDB checks the equality against { country: "US" } *as an object, not as a pattern!*

More on nested structures

- Dot notation for XPath-like path expressions
 - Books where some subsection title contains “1.1”
`db.bib.find({ "sections.sections.title": /1\.1/ })`
 - Note we that need to quote the expression
 - Again, if returns multiple things, the condition only needs to hold for at least one of them
- Use `$elemMatch` to ensure that the same array element satisfies multiple conditions, e.g.:
`db.bib.find({ sections: { $elemMatch: {
 title: /Section/, "sections.title": /1\.1/
}}})`
- Dot notation for specifying array elements
 - Books whose first author is Abiteboul
`db.bib.find({ "authors.0": "Abiteboul" })`
 - Note 0-based indexing; again, need to quote the expression

find() with projection and sorting

- List just the book prices and nothing else

```
db.bib.find({ price: { $exists: true } },  
           { _id: false, price: true })
```

 - The (optional) second argument to find() specifies projection: true means to return, false means to omit
 - `_id` is returned by default unless otherwise specified
- List books but not subsections, ordered by ISBN

```
db.bib.find({}, {"sections.sections": false}).sort({ISBN:1})
```

 - Output from find() is further sorted by sort(), where 1/-1 mean ascending/descending order
- “Aggregation pipelines” (next) are better suited for constructing more complex output

MongoDB aggregation pipeline

- Idea: think of a query as performing a sequence of “stages,” each transforming an input sequence of JSON objects to an output sequence of JSON objects
- “Aggregation” is a misnomer: there are all kinds of stages
 - Selection (`$match`), projection (`$project`), sorting (`$sort`)
 - Much of which `find()` and `sort()` already do
 - Computing/adding attributes with generalized projection (`$project/$addField`), unnesting embedded arrays (`$unwind`), and restructuring output (`$replaceRoot`)
 - Operators to transform/filter arrays (`$map/$filter`)
 - Join (`$lookup`)
 - Grouping and aggregation (`$group`)
 - Operators to aggregate (e.g., `$sum`) or collect into an array (`$push`)

The congress MongoDB database

- As in your A4
- Two collections, `people` and `committees`
 - Each object in `people` is a legislator
 - `roles` = array of objects
 - Each object in `committees` is a committee
 - `members` = array of objects
 - `subcommittees` = an array of subcommittee objects, each with its own `members` array
 - Each member object's `id` field references a legislator `_id`

`_id` is the key, and `id` is the foreign key of the `people` collection?
Ask professor to clarify about this

```
[
  {
    "_id" : "B000944",
    "birthday" : ISODate("1952-11-09T00:00:00Z"),
    "gender" : "M",
    "name" : "Sherrod Brown",
    "roles" : [
      {
        "district" : 13,
        "enddate" : ISODate("1995-01-03T00:00:00Z"),
        "party" : "Democrat",
        "startdate" : ISODate("1993-01-05T00:00:00Z"),
        "state" : "OH",
        "type" : "rep"
      },
      {
        "district" : 13,
        "enddate" : ISODate("1997-01-03T00:00:00Z"),
        "party" : "Democrat",
        "startdate" : ISODate("1995-01-04T00:00:00Z"),
        "state" : "OH",
        "type" : "rep"
      }
    ], ...
  },
  ...
]
```

```
[
  {
    "_id" : "HSAG",
    "displayname" : "House Committee on Agriculture",
    "type" : "house",
    "members" : [
      {
        "id" : "C001062",
        "role" : "Ranking Member"
      },
      {
        "id" : "T000467"
      },
      ...
    ],
    "subcommittees" : [
      {
        "code" : "15",
        "displayname" : "Conservation and Forestry",
        "members" : [
          {
            "id" : "S001209",
            "role" : "Chair"
          },
          {
            "id" : "F000455"
          },
          ...
        ]
      },
      ...
    ],
    ...
  },
  ...
]
```

Selection/projection/sorting

Find Republican legislators, output only their name and gender, sort by name

```
db.people.aggregate([
  { $match: {
    "roles.party": "Republican"
  } },
  { $project: {
    _id: false,
    name: true,
    gender: true
  } },
  { $sort: {
    name: 1
  } }
])
```

- `aggregate()` takes an array of stages
 - Hint: write/debug one at a time!
- Note again quoting the dot notation
- Note again the semantics of comparing a list of values: i.e., the query finds legislators who have ever served roles as Republicans

Generalized projection

Find Republican legislators, output their name, gender, and roles as an array of types (sen or rep)

```
db.people.aggregate([
  { $match: {
    "roles.party": "Republican"
  } },
  { $addFields: {
    compact_roles: {
      $map: { input: "$roles",
              as: "role",
              in: "$$role.type" }
    }
  } },
  { $project: {
    _id: false,
    name: true,
    gender: true,
    roles: "$compact_roles"
  } }
])
```

- Use “ : “\$xxx” ” to tell MongoDB to interpret xxx as a field in the “current” object instead of just a string literal
- In \$map, as defines a new variable to loop over elements in the input array
- For each input element, \$map computes the in expression and appends its value to the output array
 - Use “ : “\$\$xxx” ” to tell MongoDB that xxx is a new variable created during execution (as opposed to a field in the current object)

Join

For each committee (ignore its subcommittees), display its name and the name of its chairman

```
db.committees.aggregate([
  { $addFields: {
    chair_member: { $filter: {
      input: "$members",
      as: "member",
      cond: { $regexMatch: {
        input: "$$member.role",
        regex: /^(Co)?[cC]hair/
      } }
    } }
  } },
  { $lookup: {
    from: "people",
    localField: "chair_member.id",
    foreignField: "_id",
    as: "chair_person"
  } },
  { $project: {
    _id: false,
    name: "$displayname",
    chair: { $arrayElemAt: ["$chair_person.name", 0] }
  } },
])
```

- *\$filter filters input array according to cond and produces an output array*
- *In \$lookup, localField specifies the attribute in the current object whose value will be used for lookup*
- *from specifies the collection in which to look for joining objects; foreignField specifies the attribute therein to be joined*
- *\$lookup creates an attribute in the current object with the name specified by as, and sets its value to an array holding all joining objects*
- ☞ *Non-equality joins are also possible, with more complex syntax*

*\$arrayElemAt extracts an array element by its index
("chair_person.0.name" doesn't work here)*

Unnesting and restructuring

Create a list of subcommittees: for each, simply display its name, its members, and the id/name of the parent committee

```
db.committees.aggregate([
  { $unwind: "$subcommittees" },
  { $replaceRoot: { newRoot: {
    name: "$subcommittees.displayname",
    members: "$subcommittees.members",
    parent_committee: { id: "$_id",
                        name: "$displayname" }
  } } }
])
```

For each input committee, \$unwind loops over its subcommittees array, one element at a time, and outputs a copy of the committee, with its subcommittees value replaced with this single element

- *By default, \$unwind ignores committees with no subcommittees, but there is an option to keep them, with subcommittees set to null*

Grouping and aggregation

- Count legislators by gender, and list the names of legislators for each gender

```
db.people.aggregate([
  { $group: {
    _id: "$gender",
    count: { $sum: 1 },
    list: { $push: "$name" }
  }
}]
```

- *The required `_id` specifies the grouping expression, whose value becomes the identifying attribute of output objects (one per group)*
- *Other attributes hold aggregate values, computed using “accumulator” operators*
 - `$sum` compute a total by adding each input
 - `$push` creates an array by appending each input
- *Array-producing accumulator operators allows “nesting”*

Array operators vs. unnest/nest

Don't array operators \$map/\$filter look like projection/selection to you?

- You can always unnest, project/select, and then nest (aggregate) them back!
- In **nested relational algebra**, which could serve as the theoretical foundation for querying nested data, you just need the following operators:
 - $\cup, \cap, -, \sigma_p, \pi_L$
 - tup_create/destroy (which enable \times and ρ_{\dots})
 - set_create/destroy (which further enable **unnest**)
 - **nest**

(Interestingly, a more expressive language can be obtained by replacing nest by powerset, which would allow you to express transitive closure!)

Example of array ops vs. unnest/nest

Find Republican legislators, output their name, gender, and roles as an array of types (sen or rep)

```
db.people.aggregate([
  { $match: {
    "roles.party": "Republican"
  } },
  { $addFields: {
    compact_roles: {
      $map: { input: "$roles",
        as: "role",
        in: "$$role.type" }
    }
  } },
  { $project: {
    _id: false,
    name: true,
    gender: true,
    roles: "$compact_roles"
  } }
])
```

```
db.people.aggregate([
  { $unwind: "$roles" },
  { $match: {
    "roles.party": "Republican"
  } },
  { $addFields: {
    compact_role: "$roles.type"
  } },
  { $group: {
    _id: "$_id",
    name: { $last: "$name" },
    gender: { $last: "$gender" },
    compact_roles: { $push:
      "$compact_role" }
  } },
  { $project: { _id: false } }
])
```

Summary and discussion

- JSON is like a lightweight version of XML
 - But perhaps not as good for mixed contents
- Writing MongoDB queries in JSON format is sometimes convenient, but confusing in many situations
- **Query as as pipeline \approx algebra**: less “declarative,” but arguably easier to implement (especially to parallelize)
- Nested structures require more query constructs
 - They really just boil down to some form of **unnest** and **nest**: `$unwind` and `$group` in MongoDB
 - `$elemMatch/$map/$filter/$arrayElemAt` are just syntactic sugar

☞ Alternatives to MongoDB

- N1SQL: SQL-like language for JSON by CouchDB; very clean design
- JSONiq (lesser known): XQuery-like language for JSON